

# Week 1 Homework

## Question 2.1

Describe a situation or problem from your job, everyday life, current events, etc., for which a classification model would be appropriate. List some (up to 5) predictors that you might use.

**Answer:**

In Petroleum Engineering, defining whether a rock section intersected by a well is a hydrocarbon bearing reservoir or not and whether it can be economically produced is critical for the future utilization of the well.

Pay is an expression that denotes a portion of a reservoir that contains economically recoverable hydrocarbons.

In order for a rock to be considered a pay rock, It must contain sufficient:

1. Porosity (void space in the reservoir rock to store the hydrocarbons),
2. Hydrocarbon saturation as opposed to water saturation,
3. Permeability to transmit the hydrocarbons through the reservoir and to the wellbore,
4. Reserves to be economically developed, etc.

The definition of sufficient varies from geological zone to another. Also the definition of economic reserves varies significantly depending on the location (onshore vs offshore) and many other factors.

Reference: <https://glossary.oilfield.slb.com/en/Terms/p/pay.aspx>

## Question 2.2

The files `credit_card_data.txt` (without headers) and `credit_card_data-headers.txt` (with headers) contain a dataset with 654 data points, 6 continuous and 4 binary predictor variables. It has anonymized credit card applications with a binary response variable (last column) indicating if the application was positive or negative. The dataset is the "Credit Approval Data Set" from the UCI Machine Learning Repository (<https://archive.ics.uci.edu/ml/datasets/Credit+Approval>) without the categorical variables and without data points that have missing values.

### Part 1

Using the support vector machine function `ksvm` contained in the R package `kernlab`, find a good classifier for this data. Show the equation of your classifier, and how well it classifies the data points in the full data set.

### Answer Part 1

```
In [1]: # loading the dataset
# READ DATASET as Matrix
data <- as.matrix(read.table("credit_card_data-headers.txt", header = TRUE, sep = "\t"))
# Display Data
head(data)
```

A1	A2	A3	A8	A9	A10	A11	A12	A14	A15	R1
1	30.83	0.000	1.25	1	0	1	1	202	0	1
0	58.67	4.460	3.04	1	0	6	1	43	560	1
0	24.50	0.500	1.50	1	1	0	1	280	824	1
1	27.83	1.540	3.75	1	0	5	0	100	3	1
1	20.17	5.625	1.71	1	1	0	1	120	0	1
1	32.08	4.000	2.50	1	1	0	0	360	0	1

```
In [2]: # convert data to data frame for easier handling and plotting using ggplot2 library
data_df <- as.data.frame(data)
```

## Describe data properties

```
In [3]: print("Number of data points , columns in data")
        dim(data_df)
        print("Data Columns summary")
        summary(data_df)
```

```
[1] "Number of data points , columns in data"
```

1. 654

2. 11

```
[1] "Data Columns summary"
```

A1		A2		A3		A8	
Min.	:0.0000	Min.	:13.75	Min.	: 0.000	Min.	: 0.000
1st Qu.:	0.0000	1st Qu.:	22.58	1st Qu.:	1.040	1st Qu.:	0.165
Median :	1.0000	Median :	28.46	Median :	2.855	Median :	1.000
Mean :	0.6896	Mean :	31.58	Mean :	4.831	Mean :	2.242
3rd Qu.:	1.0000	3rd Qu.:	38.25	3rd Qu.:	7.438	3rd Qu.:	2.615
Max.	:1.0000	Max.	:80.25	Max.	:28.000	Max.	:28.500

A9		A10		A11		A12	
Min.	:0.0000	Min.	:0.0000	Min.	: 0.000	Min.	:0.0000
1st Qu.:	0.0000	1st Qu.:	0.0000	1st Qu.:	0.000	1st Qu.:	0.0000
Median :	1.0000	Median :	1.0000	Median :	0.000	Median :	1.0000
Mean :	0.5352	Mean :	0.5612	Mean :	2.498	Mean :	0.5382
3rd Qu.:	1.0000	3rd Qu.:	1.0000	3rd Qu.:	3.000	3rd Qu.:	1.0000
Max.	:1.0000	Max.	:1.0000	Max.	:67.000	Max.	:1.0000

A14		A15		R1	
Min.	: 0.00	Min.	: 0	Min.	:0.0000
1st Qu.:	70.75	1st Qu.:	0	1st Qu.:	0.0000
Median :	160.00	Median :	5	Median :	0.0000
Mean :	180.08	Mean :	1013	Mean :	0.4526
3rd Qu.:	271.00	3rd Qu.:	399	3rd Qu.:	1.0000
Max.	:2000.00	Max.	:100000	Max.	:1.0000

### summary of data

As described in the assignment,

1. the data contains 654 data points
2. 6 continuous variables (A2, A3, A8, A11, A14, A15)
3. 4 binary variables (A1, A9, A10, A12)
4. R1 the binary response column

for the continous columns, for example

A2 ranges from 13.75 to 80.25

A14 ranges from 0 to 2000

A15 ranges from 0 to 100,000

it's clear that the data spans a variable wide range. As a result, scaling the data would be important to ensure accurate results

## Start of SVM Analysis

Loading the kkn "kernlab: Kernel-Based Machine Learning Lab" package

```
In [4]: # load required library
# install.packages("kernlab")
options(warn=-1) # used to suppress warnings
library(kernlab)
```

A Sensitivity analysis was done on four different models (*note for all models the data was scaled*):

1. rbfdot: Radial Basis kernel "Gaussian"
2. polydot: Polynomial kernel
3. vanilladot: Linear kernel
4. tanhdot: Hyperbolic tangent kernel

For each model, C was varied on an exponential manner from 0.0001 to 100,000 (10 Steps)

Finally Model Accuracy was calculated as follows:

Accuracy = Number of correctly predicted data points \ total number of data points

```
In [5]: # list of all models
list_models <- c("polydot", "vanilladot", "rbfdot", "tanhdot")
model_names <- c("Polynomial Model", "Linear Model", "Gaussian Model", "Hyperbolic tange

# Sensitivity on C values
C_list <- c(0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000, 100000)

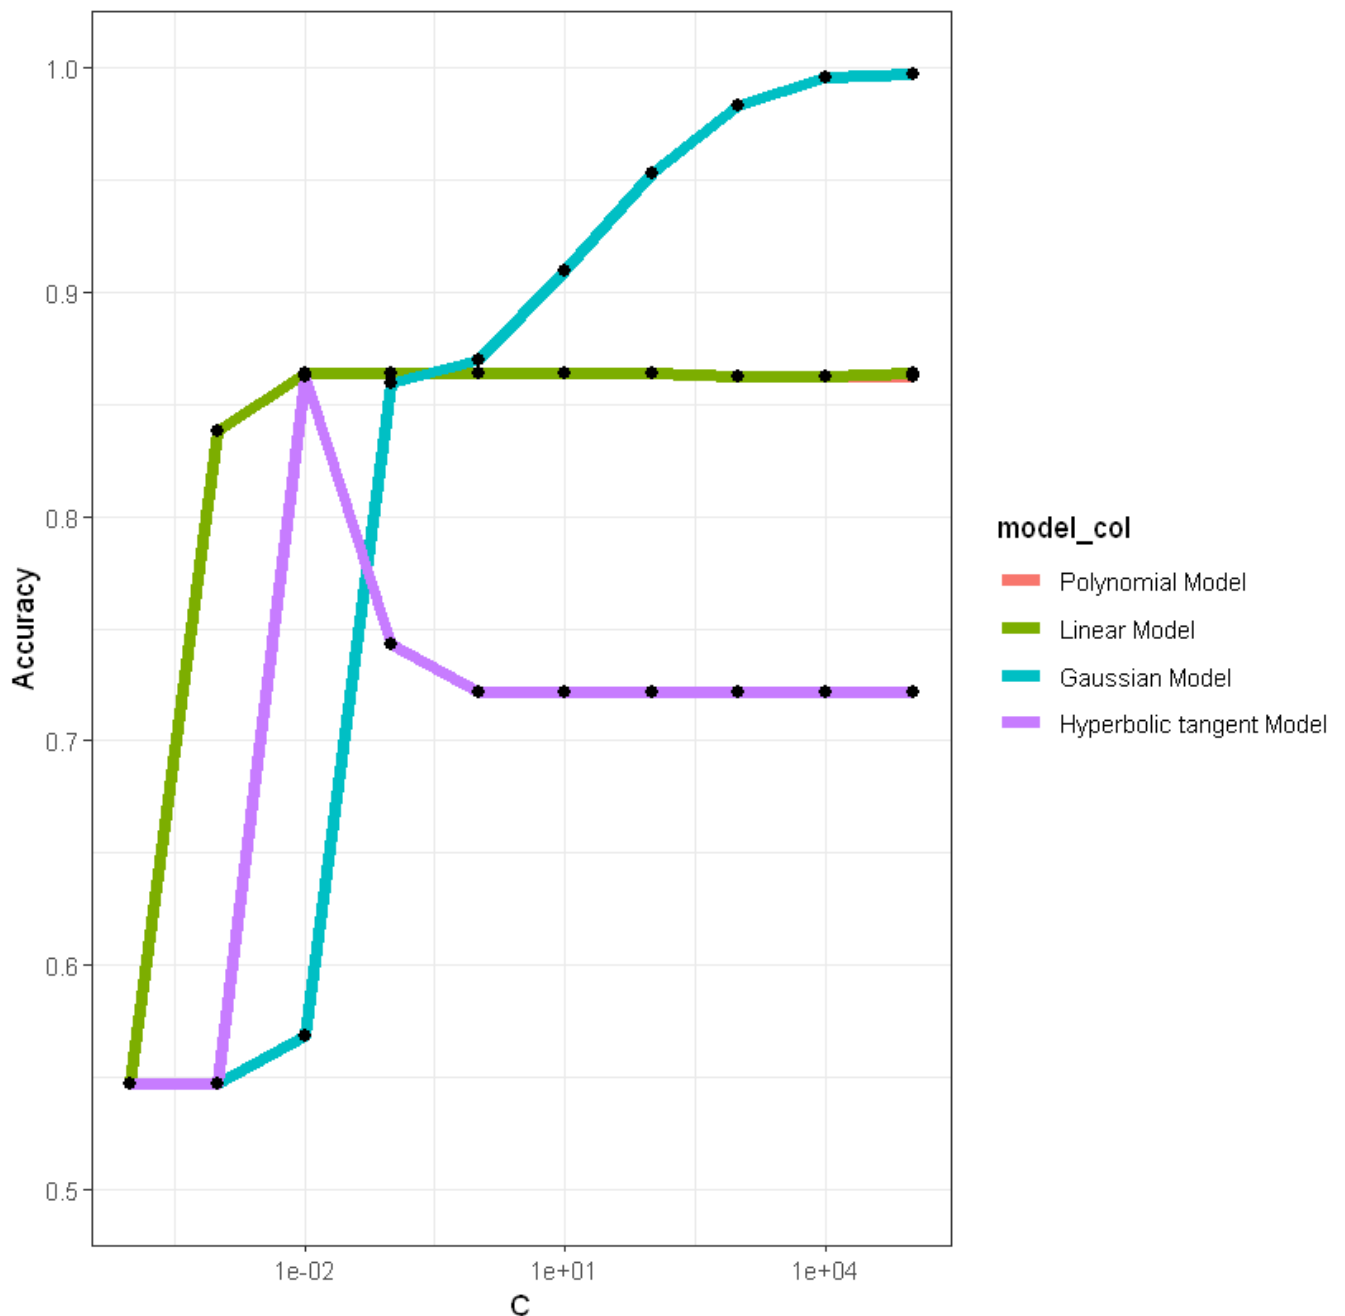
# vector to store Accuracy Results
Accuracy <- seq(1,10,1)

# DataFrame to store all the results
results_df <- data.frame()
```

```
In [6]: # Outer loop for different models
for (model in 1:4)
{
  # Inner loop for different C Values
  for (i in 1:10)
  {
    # Create Model
    model_svm <- ksvm(data[,1:10], data[,11], type="C-svc", kernel=list_models[model], C=C_list[i])
    # see what the model predicts
    pred <- predict(model_svm,data[,1:10])
    # see what fraction of the model's predictions match the actual classification (Accuracy)
    Accuracy[i] <- sum(pred == data[,11]) / nrow(data)
  }
  # Creates a vector of model name
  model_col <- rep(model_names[model], times=10)
  # Binds the results to the results dataframe
  results_df <- rbind(results_df, data.frame(C_list, Accuracy, model_col))
}
```



## Impact of C on Accuracy of Different Models



Note the X axis was transformed to log scale to better show the results given the wide range of C values used

It was noted that when  $C = 0.0001$ , the model accuracy was at it's worst. As a result, The  $C = 0.0001$  model was tested to see the frequency of 1 and 0 responses.

```
In [8]: # Create C=0.0001 Model
poor_C_model_svm <- ksvm(data[,1:10], data[,11], type="C-svc", kernel="vanilladot", C=0.0001)
# see what the model predicts
pred <- predict(poor_C_model_svm,data[,1:10])
# see Model's Accuracy
poor_C_model_svm_Accuracy <- round(sum(pred == data[,11]) / nrow(data),4)
print("Poor C model accuracy")
poor_C_model_svm_Accuracy
print("Number of False responses R1==0 predicted by the model")
sum(pred == 0)
print("Percentage of False responses R1==0 predicted by the model")
sum(pred == 0)/nrow(data)*100
```

```
Setting default kernel parameters
[1] "Poor C model accuracy"
```

0.5474

```
[1] "Number of False responses R1==0 predicted by the model"
```

654

```
[1] "Percentage of False responses R1==0 predicted by the model"
```

100

As a result, As highlighted in the Assignment description, when C is poorly selected, the majority of the responses will be either 0 or 1 (in our case, when C=0.0001, all responses will be 0)

### **Conclusions from the plot**

1. Linear and Polynomial models are overlapping (producing the same results for the same C value)
2. For both Linear and Polynomial models, the Accuracy increases with increasing C value up to C = 0.01 then flattens at 0.862 for any higher values of C. Further models will use a C = 1 assumption
3. For the Hyperbolic tangent model, the optimal C value is 0.01 with any further increase in C value results in a reduction in the Accuracy. Note that if that model was to be used, another loop to refine the C value would be needed given the wide range of the C steps.
4. For the Gaussian model, increasing the C value increases the accuracy of the model. However with C = 100,000, the model's accuracy increases to ~1.0. Further validation and testing would be required as the model could be over-fitting.

## Creating the basic Linear model for further analyses

```
In [9]: # Create Basic Model
model_svm <- ksvm(data[,1:10], data[,11], type="C-svc", kernel="vanilladot", C=1, scaled=TRUE)
# see what the model predicts
pred <- predict(model_svm,data[,1:10])
# see Model's Accuracy
basic_model_accuracy <- round(sum(pred == data[,11]) / nrow(data),4)
print("Basic Linear model accuracy")
basic_model_accuracy
```

```
Setting default kernel parameters
[1] "Basic Linear model accuracy"
```

0.8639

Quality checking the fraction of True responses ( $R1 == 1$ ) in both the data set and the prediction

```
In [10]: cat("Fraction of True responses in Data set", round(sum(data[, "R1"]==1)/nrow(data),4))
```

Fraction of True responses in Data set 0.4526

```
In [11]: cat("Fraction of True responses in prediction", round(sum(pred==1)/nrow(data),4))
```

Fraction of True responses in prediction 0.5367

Conclusion The model appears to overestimate the True responses however not excessively.

```
In [12]: # calculate a0
a0 <- model_svm@b
print("a0, the intercept is")
round(a0,4)
# calculate a1...am
a <- colSums(model_svm@xmatrix[[1]] * model_svm@coef[[1]])
print("The cofficeints of the different columns")
round(a,4)
```

```
[1] "a0, the intercept is"
```

-0.0815

```
[1] "The cofficeints of the different columns"
```

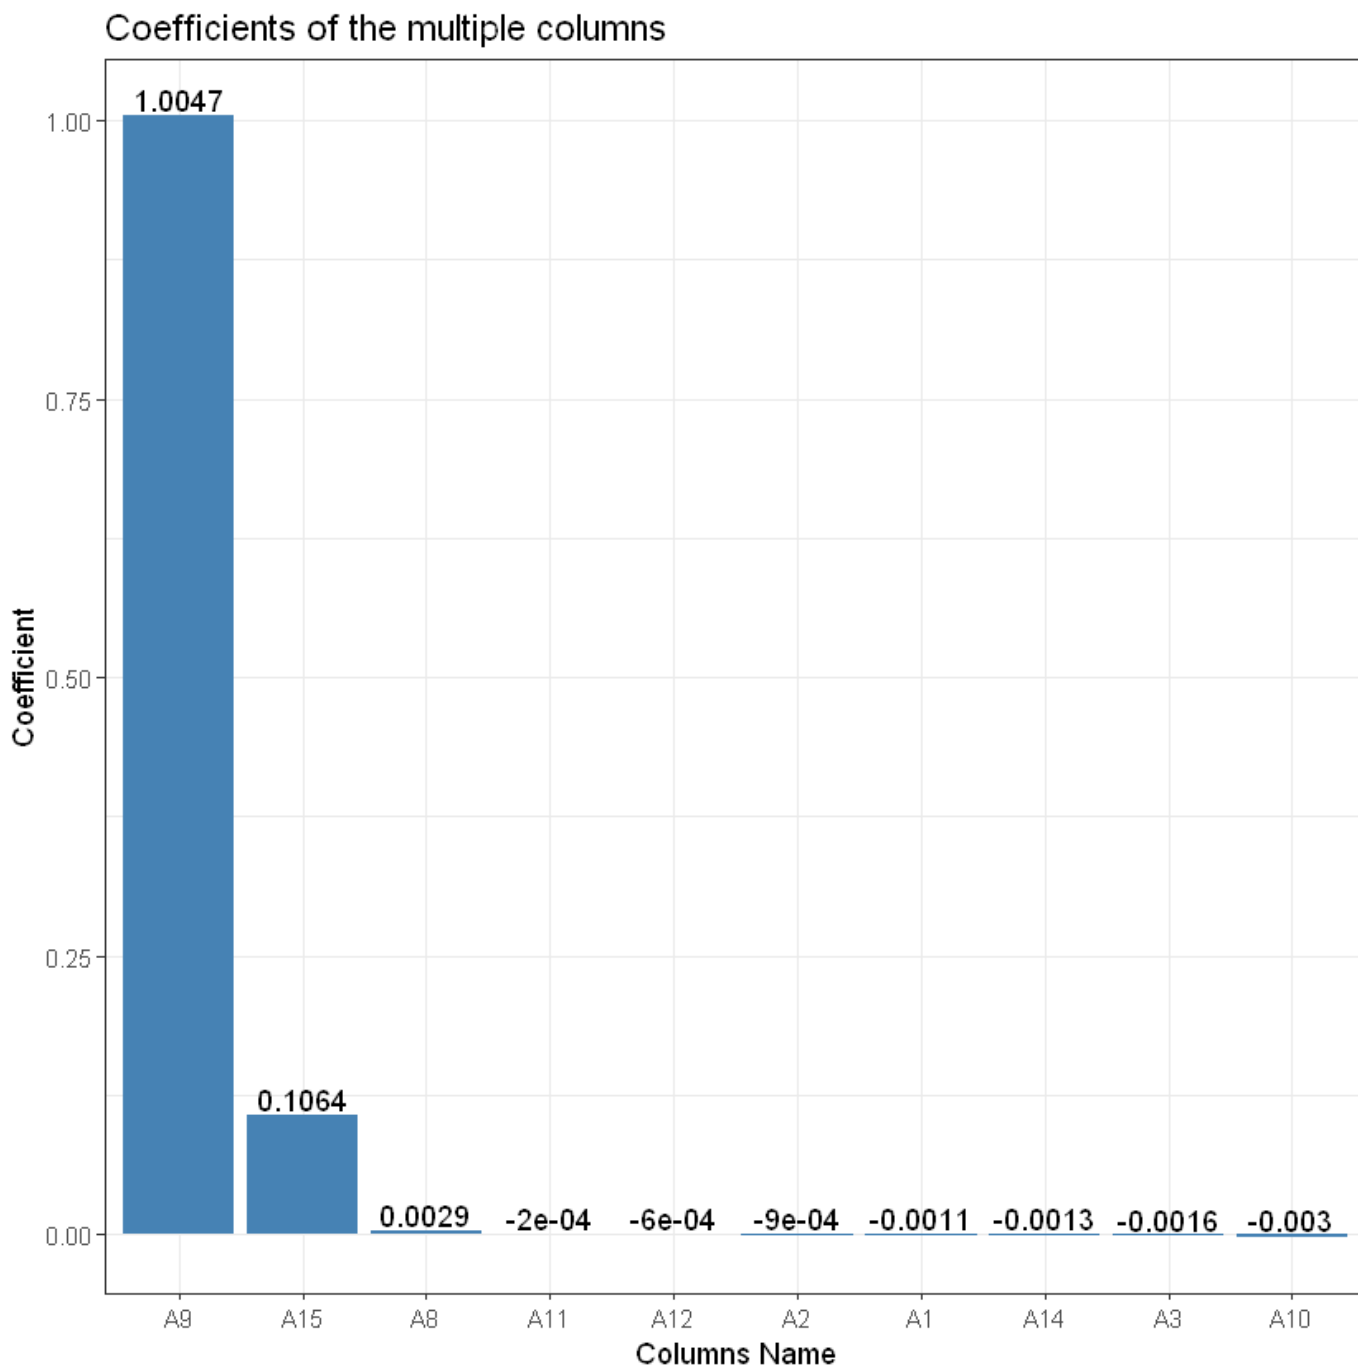
A1	-0.0011
A2	-9e-04
A3	-0.0016
A8	0.0029
A9	1.0047
A10	-0.003
A11	-2e-04
A12	-6e-04
A14	-0.0013
A15	0.1064



## Analyzing the Cofccients results

In [13]:

```
# plot coefficients
# create a dataframe to store the cofficients
a_df <- data.frame(names(a), a)
colnames(a_df) <- c("Column_name", "Coefficient")
# Plot the data
Coff_plot <- ggplot(a_df, aes(x=reorder(Column_name, -Coefficient), y=Coefficient, group=
  geom_bar(stat="identity", fill="steelblue")+
  labs(x="Columns Name", Y="Coefficient")+
  ggtitle("Coefficients of the multiple columns")+
  geom_text(aes(label=round(Coefficient,4)), position=position_dodge(width=0.9), vjust=-
  theme_bw()
print(Coff_plot)
```



### Conclusion from the plot

2 Columns (A9 & A15) has a high coffcient indicating a strong correlation with the response vector while all other columns' coffceints are almost 0

To Validate the concept, A simple model was built using only these two columns. The concept is to compare the simple model accuracy with the initial model accuracy

```
In [14]: # Building Simple Model
simple_model_svm <- ksvm(data[,c("A9", "A15")], data[,11], type="C-svc", kernel="vanilla")
# see what the model predicts
simple_pred <- predict(simple_model_svm, data[,c("A9", "A15")])
# see what fraction of the model's predictions match the actual classification
simple_model_accuracy <- round(sum(simple_pred == data[,11]) / nrow(data), 4)
print("Simple Linear model accuracy")
simple_model_accuracy
print("Remember Initial Linear model accuracy")
basic_model_accuracy
```

Setting default kernel parameters

```
[1] "Simple Linear model accuracy"
```

0.8639

```
[1] "Remember Initial Linear model accuracy"
```

0.8639

Conclusion, Both models results the in same accuracy. Only Column A9 and A15 are correlatble with the R1 response

```
In [15]: # Calculate cofficients of the simple model
# calculate a0
a0 <- simple_model_svm@b
print("a0, the intercept is")
round(a0, 4)
# calculate a1...am
a <- colSums(simple_model_svm@xmatrix[[1]] * simple_model_svm@coef[[1]])
print("The cofficeints of the different columns")
round(a, 4)
```

```
[1] "a0, the intercept is"
```

-0.0813

```
[1] "The cofficeints of the different columns"
```

A9	1.0086
----	--------

A15	0.106
-----	-------

## Final Models equations

### Initial Model equation

$$\text{Response} = -0.0011 \times A1 - 9e-04 \times A2 - 0.0016 \times A3 + 0.0029 \times A8 + 1.0047 \times A9 - 0.003 \times A10 - 2e-04 \times A11 - 6e-04 \times A12 - 0.0013 \times A13 + 0.1064 \times A15 - 0.0813$$

### Final Model equation

$$\text{Response} = 1.0086 \times A9 + 0.106 \times A15 - 0.0813$$

For both models,

1. if Response  $\geq 0$  then Response = 1
2. if Response  $< 0$  then Response = 0

### Manual calculations using the equation for the simple model

In [16]:

```
# rescaling columns
temp_data <- data
temp_data[, "A15"] <- scale(data[, "A15"])
temp_data[, "A9"] <- scale(data[, "A9"])
# Manual Calculations
Response <- (temp_data[, "A9"] * 1.0086 + temp_data[, "A15"] * 0.106 - 0.0813)
# If Response >= 0 , 1 else 0
Response_final <- as.numeric(Response >= 0)
# print total number of mis-matches
print("Total number of mismatches")
sum(Response_final != simple_pred)
```

```
[1] "Total number of mismatches"
```

```
0
```

In conclusion, The KSVM model and the manual equation are identical

As a final Quality and logic check since A9 column coefficient is  $\sim 1$ , A9 (which is a binary column) should be directly correlated to the response vector. Using it alone to predict the response vector should result in a high accuracy

In [17]:

```
# print Correlation for column A9 alone
A9_Accuracy <- sum(data[, "A9"] == data[, "R1"]) / nrow(data)
print(c("A9 Correlation with Target column (R1)", round(A9_Accuracy, 4)))
```

```
[1] "A9 Correlation with Target column (R1) "
```

```
[2] "0.8624"
```

Indeed using A9 as a linear direct predictor results in an accuracy of 0.8624 as compared to 0.8639 for the model including A15 or all other columns

## **Part 1 Conclusions:**

- 1. Based on testing, two columns only A9 & A15 has the highest impact on the response.**
- 2. Using these 2 columns (or all the columns) has an accuracy of 86.4%**
- 3. Any C value  $\geq 0.01$  and  $\leq 100,000$  (limit of testing) results in the same Linear SVM Model results.**
- 4. The model equation is**

$$\text{Response} = 1.0086 \times A9 + 0.106 \times A15 - 0.0813$$

- 1. if Response  $\geq 0$  then Response = 1**
- 2. If Response  $< 0$  then Response = 0**

## Part 2

Using the k-nearest-neighbors classification function `kkn` contained in the R `kkn` package, suggest a good value of `k`, and show how well it classifies that data points in the full data set. Don't forget to scale the data (`scale=TRUE` in `kkn`)

```
In [18]: # load required library
#install.packages("kkn")
options(warn=-1) # used to suppress warnings
library(kkn)
```

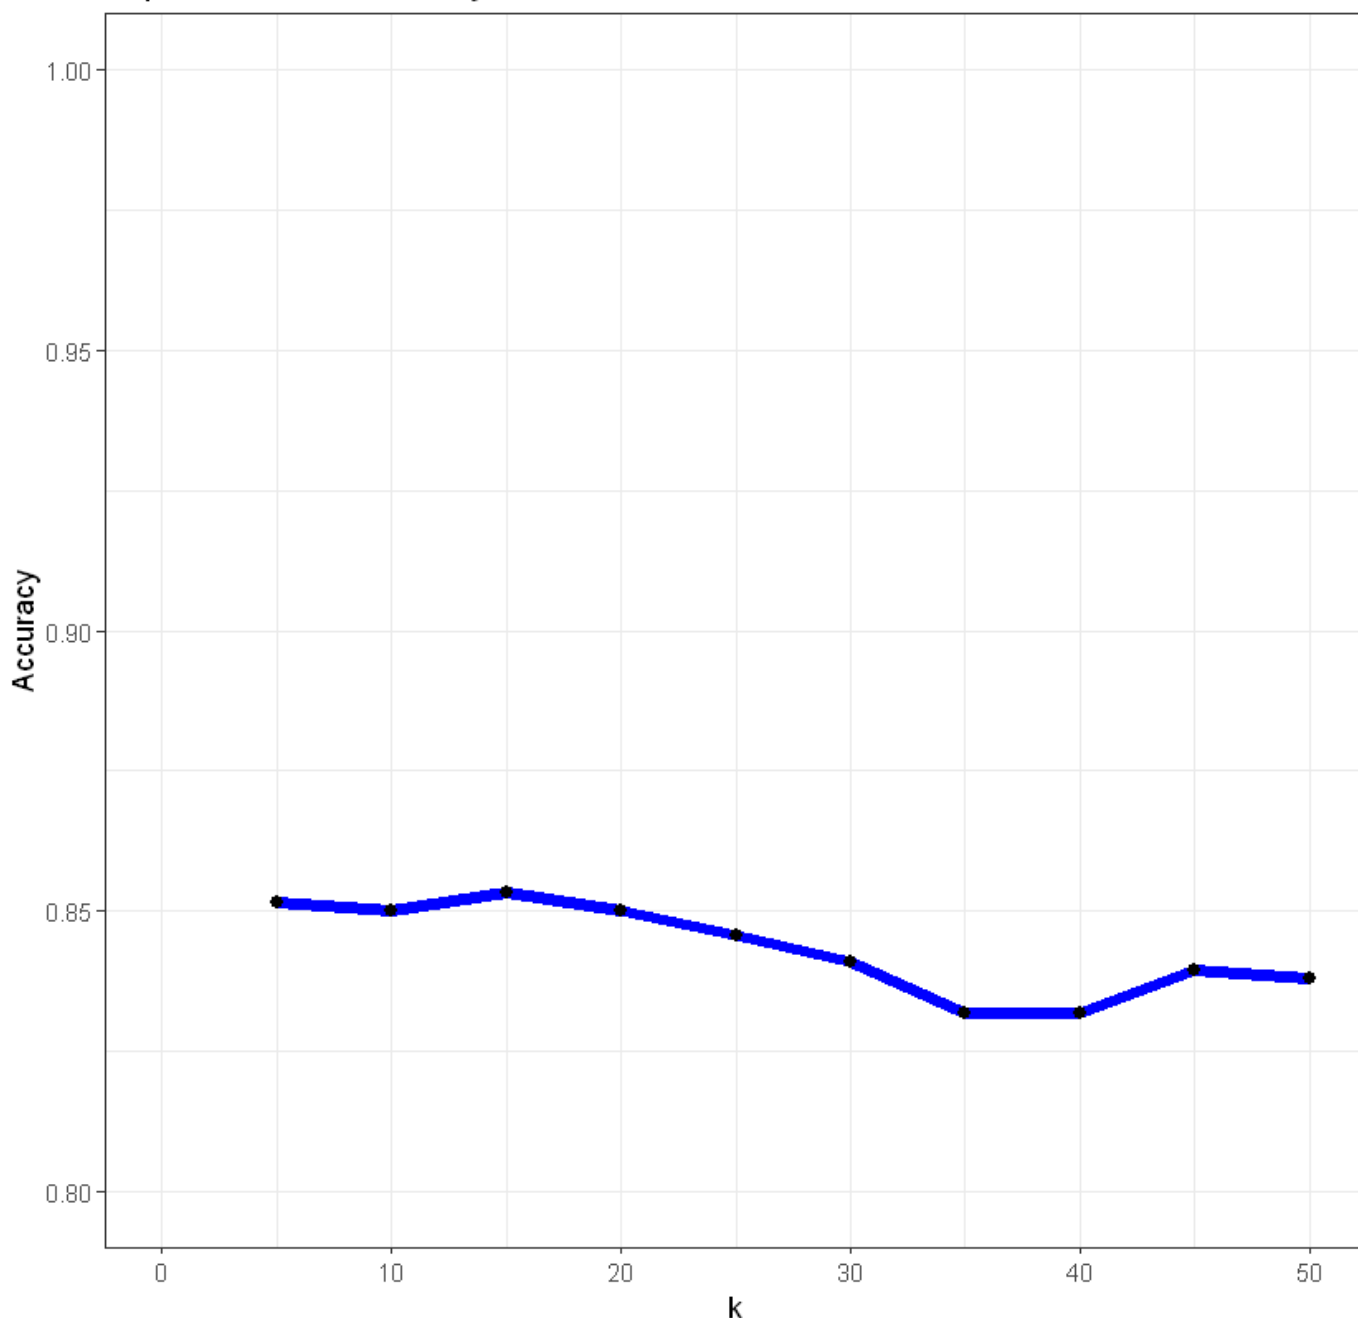
An initial Sensitivity analysis was done for `k` values ranging from 5 to 50 increment 5 Model accuracy was calculated as defined previously

```
In [19]: # KNN Solution
Accuracy <- seq(1,10,1)
k_list <- seq(5,50,5)
# outer k values loop
for (j in 1:10)
{
  N <- nrow(data_df)
  results <- vector(length=N)
  # inner loop for all rows
  for (i in 1:N) {
    KNN_model <- kkn(data_df[-i,11]~., train=data_df[-i,1:10], test=data_df[i,1:10], k=
    results[i] <- fitted(KNN_model)
  }
  Accuracy[j] <- sum(round(results,0) == data[,11]) / nrow(data)
}
```

Plot the results where X-axis is `K` value and Y-axis is the model accuracy

```
In [20]: knn_df <- data.frame(k_list, Accuracy)
knn_plot <- ggplot(knn_df, aes(x=k_list, y=Accuracy)) +
  geom_line(size=2, color="blue")+
  geom_point(size=2, color="black")+
  ggtitle("Impact of k on Accuracy of KNN Model")+
  labs(x="k", Y="Accuracy")+
  ylim(0.8, 1)+
  xlim(0, 50)+
  theme_bw()
knn_plot
```

Impact of k on Accuracy of KNN Model

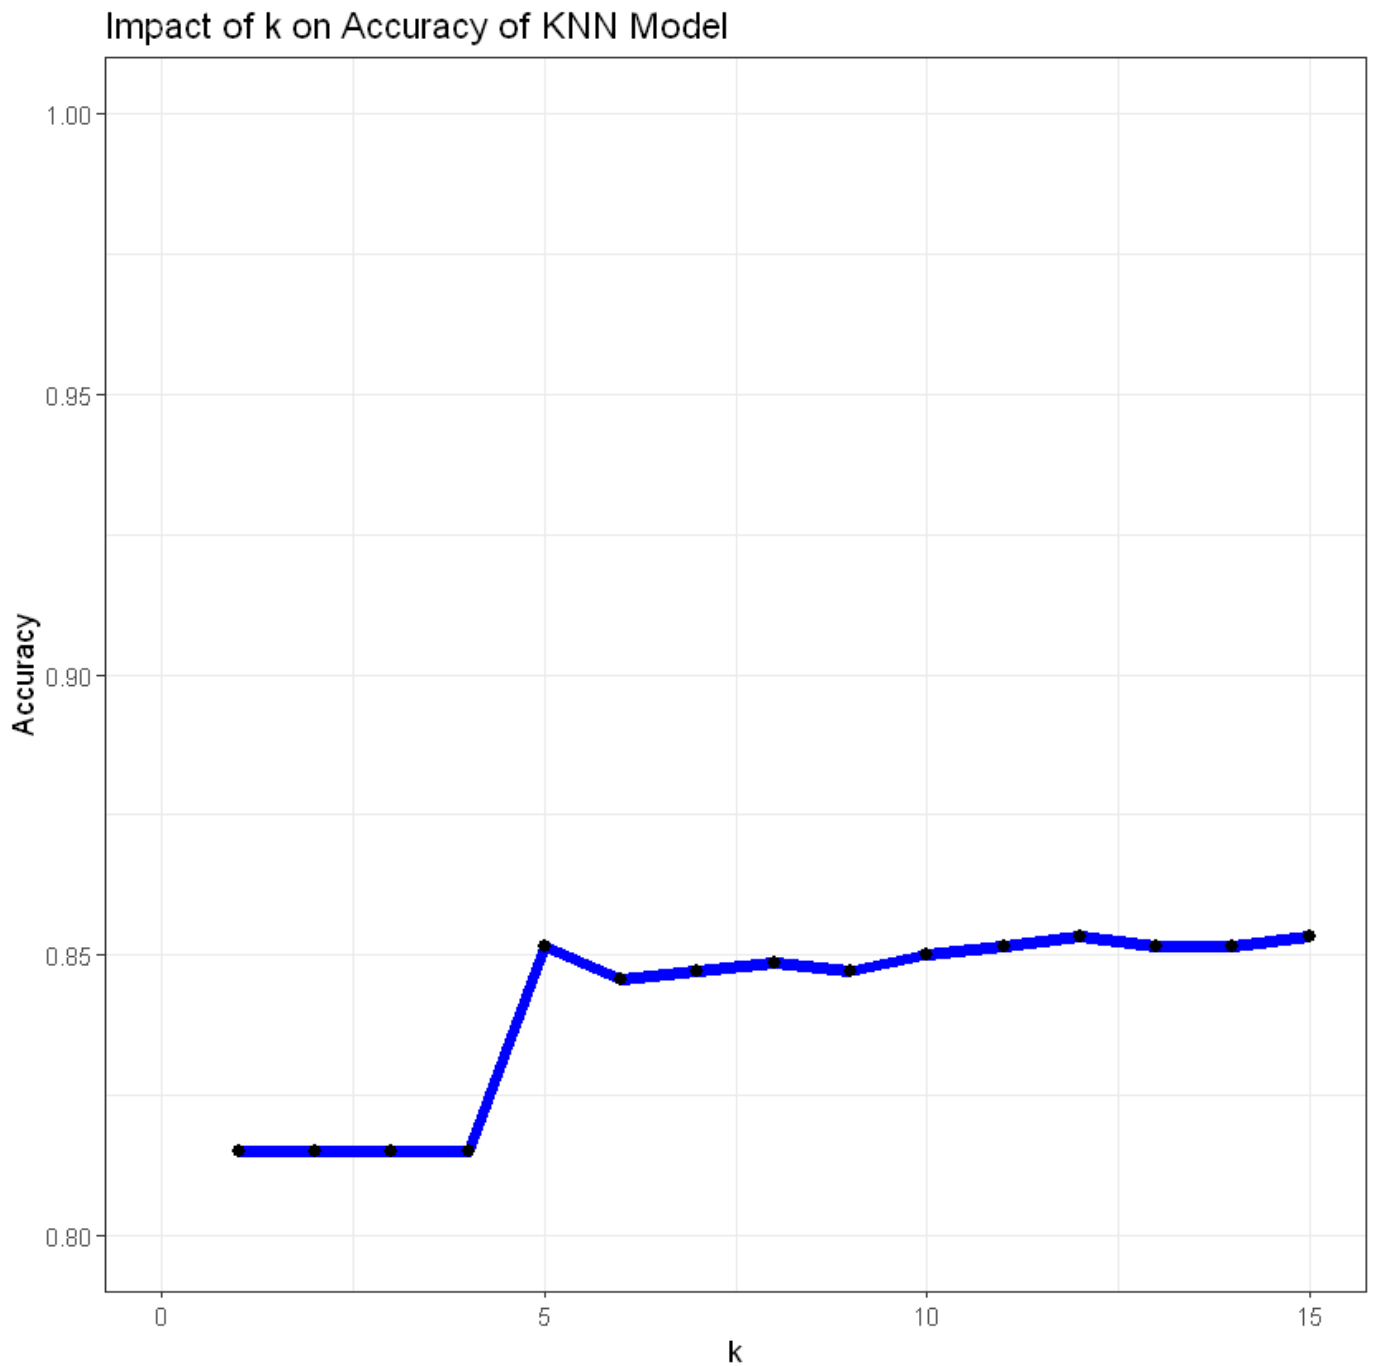


The analysis showed that the optimal K value is around 10 with further increasing or decreasing the k value results in lower Model's accuracy As a result, the experiment was repeated with k ranging from 1 to 15 with increments of 1.

```
In [21]: Accuracy <- seq(1,15,1)
k_list <- seq(1,15,1)
for (j in 1:15)
{
  N <- nrow(data_df)
  results <- vector(length=N)
  for (i in 1:N) {
    KNN_model <- kknk(data_df[-i,11]~., train=data_df[-i,1:10], test=data_df[i,1:10], k=
    results[i] <- fitted(KNN_model)
  }
  Accuracy[j] <- sum(round(results,0) == data[,11]) / nrow(data)
}
```

```
In [22]: knn_df <- data.frame(k_list, Accuracy)
```

```
knn_plot <- ggplot(knn_df, aes(x=k_list, y=Accuracy)) +
  geom_line(size=2, color="blue")+
  geom_point(size=2, color="black")+
  ggtitle("Impact of k on Accuracy of KNN Model")+
  labs(x="k", Y="Accuracy")+
  ylim(0.8, 1)+
  xlim(0, 15)+
  theme_bw()
knn_plot
```



Although with minor differences, k = 12 or k=15 results in the highest model Accuracy

## Final KNN Model

```
In [23]: N <- nrow(data_df)
results <- vector(length=N)
# inner loop for all rows
for (i in 1:N) {
  KNN_model <- kknk(data_df[-i,11]~., train=data_df[-i,1:10], test=data_df[i,1:10], k=12,
  results[i] <- fitted(KNN_model)
}
Accuracy <- sum(round(results,0) == data[,11]) / nrow(data)
print("Final KNN Model Accuracy")
round(Accuracy, 4)

[1] "Final KNN Model Accuracy"
0.8532
```

Quality checking the fraction of True responses ( $R1 == 1$ ) in both the data set and the prediction

```
In [24]: cat("Fraction of True responses in Data set", round(sum(data[, "R1"]==1)/nrow(data),4))

Fraction of True responses in Data set 0.4526
```

```
In [25]: cat("Fraction of True responses in Data set", round(sum(round(results,0)==1)/nrow(data),4))

Fraction of True responses in Data set 0.4587
```

As opposed to the Linear SVM model, KNN model although has a slightly lower accuracy (0.8532 compared to 0.8639) but it has a more accurate percentage of True responses when compared to the dataset.

Notes of the KNN model:

Looping through all the points is inefficient in terms of computation.

Alternatively, we can split the data into training and testing database to be more efficient.



# Final Conclusions

1. KNN optimal k value is 12 or 15 with model accuracy of 85.3%
2. SVM linear model optimal C value is higher than 0.01 and less than 100,000 (limit of testing) with model accuracy of 86.4%
3. Column A9 has the highest correlation with the response vector followed by columns A15 while all other columns has a much lower impact on the model results
4. The SVM linear model tends to over-estimate the True responses (54%) while KNN model has a more accurate percentage of True responses (45.9%) which is 0.6% higher than the percentage of True responses in the dataset